Sorting Heuristics on Synthetic Workloads

R. Placeholder and S. Sample, Example Institute

Abstract

We compare three fictional sorting heuristics on synthetic workloads. The results are invented for software testing and must not be cited. Values such as 0,95 and 1.000 are written exactly as an author might, and an extractor must not 'fix' them.

1 Introduction

Document extraction tools are often evaluated on tidy inputs. This note instead mixes running text, a numbered section structure, a results table and a reference list, which is the minimum structure of most short scientific reports.

2 Method

Each heuristic was run five times on each of three workload sizes. We report the median runtime in milliseconds and the relative speed-up against the baseline.

3 Results

Heuristic	n = 1k (ms)	n = 10k (ms)	n = 100k (ms)	Speed-up
Baseline	1.8	22.4	301.0	1.00
Bucketed	1.2	14.9	187.5	1.61
Adaptive	1.4	13.1	162.2	1.86

Table 1: Median runtime over five runs (invented data).

4 Discussion

The adaptive heuristic is fastest for large inputs, while the bucketed variant wins for the smallest workload. Both claims hold only for this invented data set.

References

[1] A. Author. An Imaginary Survey of Sorting. Journal of Examples, 2024.

[2] B. Writer. Placeholder Benchmarks. Proc. Workshop on Fictional Data, 2025.